



OpenShift Container Platform 4.5

Backup and restore

Backing up and restoring your OpenShift Container Platform cluster

OpenShift Container Platform 4.5 Backup and restore

Backing up and restoring your OpenShift Container Platform cluster

Legal Notice

Copyright © 2020 Red Hat, Inc.

The text of and illustrations in this document are licensed by Red Hat under a Creative Commons Attribution–Share Alike 3.0 Unported license ("CC-BY-SA"). An explanation of CC-BY-SA is available at

<http://creativecommons.org/licenses/by-sa/3.0/>

. In accordance with CC-BY-SA, if you distribute this document or an adaptation of it, you must provide the URL for the original version.

Red Hat, as the licensor of this document, waives the right to enforce, and agrees not to assert, Section 4d of CC-BY-SA to the fullest extent permitted by applicable law.

Red Hat, Red Hat Enterprise Linux, the Shadowman logo, the Red Hat logo, JBoss, OpenShift, Fedora, the Infinity logo, and RHCE are trademarks of Red Hat, Inc., registered in the United States and other countries.

Linux® is the registered trademark of Linus Torvalds in the United States and other countries.

Java® is a registered trademark of Oracle and/or its affiliates.

XFS® is a trademark of Silicon Graphics International Corp. or its subsidiaries in the United States and/or other countries.

MySQL® is a registered trademark of MySQL AB in the United States, the European Union and other countries.

Node.js® is an official trademark of Joyent. Red Hat is not formally related to or endorsed by the official Joyent Node.js open source or commercial project.

The OpenStack® Word Mark and OpenStack logo are either registered trademarks/service marks or trademarks/service marks of the OpenStack Foundation, in the United States and other countries and are used with the OpenStack Foundation's permission. We are not affiliated with, endorsed or sponsored by the OpenStack Foundation, or the OpenStack community.

All other trademarks are the property of their respective owners.

Abstract

This document provides instructions for backing up your cluster's data and for recovering from various disaster scenarios.

Table of Contents

CHAPTER 1. BACKING UP ETCD	3
1.1. BACKING UP ETCD DATA	3
CHAPTER 2. REPLACING AN UNHEALTHY ETCD MEMBER	5
2.1. PREREQUISITES	5
2.2. IDENTIFYING AN UNHEALTHY ETCD MEMBER	5
2.3. DETERMINING THE STATE OF THE UNHEALTHY ETCD MEMBER	5
2.4. REPLACING THE UNHEALTHY ETCD MEMBER	7
2.4.1. Replacing an unhealthy etcd member whose machine is not running or whose node is not ready	7
2.4.2. Replacing an unhealthy etcd member whose etcd Pod is crashlooping	13
CHAPTER 3. SHUTTING DOWN THE CLUSTER GRACEFULLY	16
3.1. PREREQUISITES	16
3.2. SHUTTING DOWN THE CLUSTER	16
CHAPTER 4. RESTARTING THE CLUSTER GRACEFULLY	18
4.1. PREREQUISITES	18
4.2. RESTARTING THE CLUSTER	18
CHAPTER 5. DISASTER RECOVERY	21
5.1. ABOUT DISASTER RECOVERY	21
5.2. RECOVERING FROM LOST MASTER HOSTS	21
5.3. RESTORING TO A PREVIOUS CLUSTER STATE	21
5.3.1. Restoring to a previous cluster state	21
5.4. RECOVERING FROM EXPIRED CONTROL PLANE CERTIFICATES	26
5.4.1. Recovering from expired control plane certificates	26

CHAPTER 1. BACKING UP ETCD

etcd is the key-value store for OpenShift Container Platform, which persists the state of all resource objects.

Back up your cluster’s etcd data regularly and store in a secure location ideally outside the OpenShift Container Platform environment. Do not take an etcd backup before the first certificate rotation completes, which occurs 24 hours after installation, otherwise the backup will contain expired certificates. It is also recommended to take etcd backups during non-peak usage hours, as it is a blocking action.



IMPORTANT

You back up your cluster’s etcd data by performing a single invocation of the backup script on a master host. Do not take a backup for each master host.

After you have an etcd backup, you can [restore to a previous cluster state](#).

You can perform the [etcd data backup process](#) on any master host that has a running etcd instance.

1.1. BACKING UP ETCD DATA

Follow these steps to back up etcd data by creating an etcd snapshot and backing up the resources for the static Pods. This backup can be saved and used at a later time if you need to restore etcd.



IMPORTANT

Only save a backup from a single master host. Do not take a backup from each master host in the cluster.

Prerequisites

- You have access to the cluster as a user with the **cluster-admin** role.
- You have checked whether the cluster-wide proxy is enabled.

TIP

You can check whether the proxy is enabled by reviewing the output of **oc get proxy cluster -o yaml**. The proxy is enabled if the **httpProxy**, **httpsProxy**, and **noProxy** fields have values set.

Procedure

1. Start a debug session for a master node:

```
$ oc debug node/<node_name>
```

2. Change your root directory to the host:

```
sh-4.2# chroot /host
```

3. If the cluster-wide proxy is enabled, be sure that you have exported the **NO_PROXY**, **HTTP_PROXY**, and **HTTPS_PROXY** environment variables.

- Run the **cluster-backup.sh** script and pass in the location to save the backup to.

```
sh-4.4# /usr/local/bin/cluster-backup.sh /home/core/assets/backup
```

Example script output

```
1bf371f1b5a483927cd01bb593b0e12cff406eb8d7d0acf4ab079c36a0abd3f7
etcdctl version: 3.3.18
API version: 3.3
found latest kube-apiserver-pod: /etc/kubernetes/static-pod-resources/kube-apiserver-pod-7
found latest kube-controller-manager-pod: /etc/kubernetes/static-pod-resources/kube-
controller-manager-pod-8
found latest kube-scheduler-pod: /etc/kubernetes/static-pod-resources/kube-scheduler-pod-6
found latest etcd-pod: /etc/kubernetes/static-pod-resources/etcd-pod-2
Snapshot saved at /home/core/assets/backup/snapshot_2020-03-18_220218.db
snapshot db and kube resources are successfully saved to /home/core/assets/backup
```

In this example, two files are created in the **/home/core/assets/backup** directory on the master host:

- **snapshot_<datetimestamp>.db**: This file is the etcd snapshot.
- **static_kuberresources_<datetimestamp>.tar.gz**: This file contains the resources for the static Pods. If etcd encryption is enabled, it also contains the encryption keys for the etcd snapshot.



NOTE

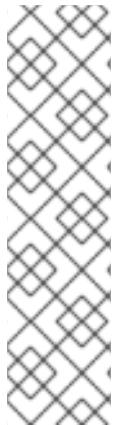
If etcd encryption is enabled, it is recommended to store this second file separately from the etcd snapshot for security reasons. However, this file is required in order to restore from the etcd snapshot.

Keep in mind that etcd encryption only encrypts values, not keys. This means that resource types, namespaces, and object names are unencrypted.

CHAPTER 2. REPLACING AN UNHEALTHY ETCD MEMBER

This document describes the process to replace a single unhealthy etcd member.

This process depends on whether the etcd member is unhealthy because the machine is not running or the node is not ready, or whether it is unhealthy because the etcd Pod is crashlooping.



NOTE

If you have lost the majority of your master hosts, leading to etcd quorum loss, then you must follow the disaster recovery procedure to [restore to a previous cluster state](#) instead of this procedure.

If the control plane certificates are not valid on the member being replaced, then you must follow the procedure to [recover from expired control plane certificates](#) instead of this procedure.

If a master node is lost and a new one is created, the etcd cluster Operator handles generating the new TLS certificates and adding the node as an etcd member.

2.1. PREREQUISITES

- Take an [etcd backup](#) prior to replacing an unhealthy etcd member.

2.2. IDENTIFYING AN UNHEALTHY ETCD MEMBER

You can identify if your cluster has an unhealthy etcd member.

Prerequisites

- Access to the cluster as a user with the **cluster-admin** role.

Procedure

1. Check the status of the **EtcdMembersAvailable** status condition using the following command:

```
$ oc get etcd -o=jsonpath='{range .items[0].status.conditions[?(@.type=="EtcdMembersAvailable")]}{.message}{"\n"}'
```

2. Review the output:

```
2 of 3 members are available, ip-10-0-131-183.ec2.internal is unhealthy
```

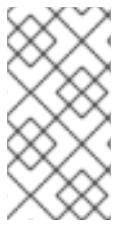
This example output shows that the **ip-10-0-131-183.ec2.internal** etcd member is unhealthy.

2.3. DETERMINING THE STATE OF THE UNHEALTHY ETCD MEMBER

The steps to replace an unhealthy etcd member depend on which of the following states your etcd member is in:

- The machine is not running or the node is not ready
- The etcd Pod is crashlooping

This procedure determines which state your etcd member is in. This enables you to know which procedure to follow to replace the unhealthy etcd member.



NOTE

If you are aware that the machine is not running or the node is not ready, but you expect it to return to a healthy state soon, then you do not need to perform a procedure to replace the etcd member. The etcd cluster Operator will automatically sync when the machine or node returns to a healthy state.

Prerequisites

- You have access to the cluster as a user with the **cluster-admin** role.
- You have identified an unhealthy etcd member.

Procedure

1. Determine if the **machine is not running**

```
$ oc get machines -A -o jsonpath='{range .items[*]}{@.status.nodeRef.name}{"\t"}{@.status.providerStatus.instanceState}{"\n"}' | grep -v running
```

Example output

```
ip-10-0-131-183.ec2.internal stopped ①
```

- ① This output lists the node and the status of the node's machine. If the status is anything other than **running**, then the **machine is not running**.

If the **machine is not running** then follow the *Replacing an unhealthy etcd member whose machine is not running or whose node is not ready* procedure.

2. Determine if the **node is not ready**.

If either of the following scenarios are true, then the **node is not ready**.

- If the machine is running, then check whether the node is unreachable:

```
$ oc get nodes -o jsonpath='{range .items[*]}{"\n"}{.metadata.name}{"\t"}{range .spec.taints[*]}{.key}{" "}' | grep unreachable
```

Example output

```
ip-10-0-131-183.ec2.internal node-role.kubernetes.io/master  
node.kubernetes.io/unreachable node.kubernetes.io/unreachable ①
```

- ① If the node is listed with an **unreachable** taint, then the **node is not ready**.

- If the node is still reachable, then check whether the node is listed as **NotReady**:

```
$ oc get nodes -l node-role.kubernetes.io/master | grep "NotReady"
```

Example output

```
ip-10-0-131-183.ec2.internal NotReady master 122m v1.18.3 1
```

- 1 If the node is listed as **NotReady**, then the **node is not ready**.

If the **node is not ready**, then follow the *Replacing an unhealthy etcd member whose machine is not running or whose node is not ready* procedure.

3. Determine if the etcd Pod is crashlooping

If the machine is running and the node is ready, then check whether the etcd Pod is crashlooping.

- a. Verify that all master nodes are listed as **Ready**:

```
$ oc get nodes -l node-role.kubernetes.io/master
```

Example output

NAME	STATUS	ROLES	AGE	VERSION
ip-10-0-131-183.ec2.internal	Ready	master	6h13m	v1.18.3
ip-10-0-164-97.ec2.internal	Ready	master	6h13m	v1.18.3
ip-10-0-154-204.ec2.internal	Ready	master	6h13m	v1.18.3

- b. Check whether the status of an etcd Pod is either **Error** or **CrashloopBackoff**:

```
$ oc get pods -n openshift-etcd | grep etcd
```

Example output

NAME	READY	STATUS	REPLICAS	AGE	VERSION
etcd-ip-10-0-131-183.ec2.internal	2/3	Error	7	6h9m	1
etcd-ip-10-0-164-97.ec2.internal	3/3	Running	0	6h6m	
etcd-ip-10-0-154-204.ec2.internal	3/3	Running	0	6h6m	

- 1 Since this status of this Pod is **Error**, then the **etcd Pod is crashlooping**.

If the **etcd Pod is crashlooping** then follow the *Replacing an unhealthy etcd member whose etcd Pod is crashlooping* procedure.

2.4. REPLACING THE UNHEALTHY ETCD MEMBER

Depending on the state of your unhealthy etcd member, use one of the following procedures:

- [Replacing an unhealthy etcd member whose machine is not running or whose node is not ready](#)
- [Replacing an unhealthy etcd member whose etcd Pod is crashlooping](#)

2.4.1. Replacing an unhealthy etcd member whose machine is not running or whose node is not ready

This procedure details the steps to replace an etcd member that is unhealthy either because the machine is not running or because the node is not ready.

Prerequisites

- You have identified the unhealthy etcd member.
- You have verified that either the machine is not running or the node is not ready.
- You have access to the cluster as a user with the **cluster-admin** role.
- You have taken an etcd backup.



IMPORTANT

It is important to take an etcd backup before performing this procedure so that your cluster can be restored if you encounter any issues.

Procedure

1. Remove the unhealthy member.

- a. Choose a Pod that is *not* on the affected node:

In a terminal that has access to the cluster as a **cluster-admin** user, run the following command:

```
$ oc get pods -n openshift-etcd | grep etcd
```

Example output

```
etcd-ip-10-0-131-183.ec2.internal      3/3  Running   0    123m
etcd-ip-10-0-164-97.ec2.internal      3/3  Running   0    123m
etcd-ip-10-0-154-204.ec2.internal     3/3  Running   0    124m
```

- b. Connect to the running etcd container, passing in the name of a Pod that is not on the affected node:

In a terminal that has access to the cluster as a **cluster-admin** user, run the following command:

```
$ oc rsh -n openshift-etcd etcd-ip-10-0-154-204.ec2.internal
```

- c. View the member list:

```
sh-4.2# etcdctl member list -w table
```

Example output

ID ADDRS	STATUS	NAME	PEER ADDRS	CLIENT

```
| 6fc1e7c9db35841d | started | ip-10-0-131-183.ec2.internal | https://10.0.131.183:2380 |
https://10.0.131.183:2379 |
| 757b6793e2408b6c | started | ip-10-0-164-97.ec2.internal | https://10.0.164.97:2380 |
https://10.0.164.97:2379 |
| ca8c2990a0aa29d1 | started | ip-10-0-154-204.ec2.internal | https://10.0.154.204:2380 |
https://10.0.154.204:2379 |
+-----+-----+-----+
-----+
```

- d. Remove the unhealthy etcd member by providing the ID to the **etcdctl member remove** command:

```
sh-4.2# etcdctl member remove 6fc1e7c9db35841d
```

Example output

```
Member 6fc1e7c9db35841d removed from cluster baa565c8919b060e
```

- e. View the member list again and verify that the member was removed:

```
sh-4.2# etcdctl member list -w table
```

Example output

ID ADDRS	STATUS	NAME	PEER ADDRS	CLIENT
757b6793e2408b6c https://10.0.164.97:2379	started	ip-10-0-164-97.ec2.internal	https://10.0.164.97:2380 https://10.0.164.97:2379	
ca8c2990a0aa29d1 https://10.0.154.204:2379	started	ip-10-0-154-204.ec2.internal	https://10.0.154.204:2380 https://10.0.154.204:2379	

You can now exit the node shell.

2. Delete and recreate the master machine. After this machine is recreated, a new revision is forced and etcd scales up automatically.

If you are running installer-provisioned infrastructure, or you used the Machine API to create your machines, follow these steps. Otherwise, you must create the new master using the same method that was used to originally create it.

- a. Obtain the machine for the unhealthy member.

In a terminal that has access to the cluster as a **cluster-admin** user, run the following command:

```
$ oc get machines -n openshift-machine-api -o wide
```

Example output

NAME NODE	PHASE PROVIDERID	TYPE	REGION	ZONE	AGE
					STATE
clustername-8qw5l-master-0	Running	m4.xlarge	us-east-1	us-east-1a	
3h37m ip-10-0-131-183.ec2.internal	aws://us-east-1a/i-0ec2782f8287dfb7e	stopped			
1 clustername-8qw5l-master-1	Running	m4.xlarge	us-east-1	us-east-1b	
3h37m ip-10-0-154-204.ec2.internal	aws://us-east-1b/i-096c349b700a19631	running			
clustername-8qw5l-master-2	Running	m4.xlarge	us-east-1	us-east-1c	
3h37m ip-10-0-164-97.ec2.internal	aws://us-east-1c/i-02626f1dba9ed5bba	running			
clustername-8qw5l-worker-us-east-1a-wbtgd	Running	m4.large	us-east-1	us-east-1a	
3h28m ip-10-0-129-226.ec2.internal	aws://us-east-1a/i-010ef6279b4662ced	running			
clustername-8qw5l-worker-us-east-1b-lrdxb	Running	m4.large	us-east-1	us-east-1b	
3h28m ip-10-0-144-248.ec2.internal	aws://us-east-1b/i-0cb45ac45a166173b	running			
clustername-8qw5l-worker-us-east-1c-pkg26	Running	m4.large	us-east-1	us-east-1c	
3h28m ip-10-0-170-181.ec2.internal	aws://us-east-1c/i-06861c00007751b0a	running			

- 1** This is the master machine for the unhealthy node, **ip-10-0-131-183.ec2.internal**.

- b. Save the machine configuration to a file on your file system:

```
$ oc get machine clustername-8qw5l-master-0 \ 1
-n openshift-machine-api \
-o yaml \
> new-master-machine.yaml
```

- 1** Specify the name of the master machine for the unhealthy node.

- c. Edit the **new-master-machine.yaml** file that was created in the previous step.

- i. Remove the entire **status** section:

```
status:
addresses:
- address: 10.0.131.183
  type: InternalIP
- address: ip-10-0-131-183.ec2.internal
  type: InternalDNS
- address: ip-10-0-131-183.ec2.internal
  type: Hostname
lastUpdated: "2020-04-20T17:44:29Z"
nodeRef:
  kind: Node
  name: ip-10-0-131-183.ec2.internal
  uid: acca4411-af0d-4387-b73e-52b2484295ad
phase: Running
providerStatus:
  apiVersion: awsproviderconfig.openshift.io/v1beta1
  conditions:
  - lastProbeTime: "2020-04-20T16:53:50Z"
    lastTransitionTime: "2020-04-20T16:53:50Z"
    message: machine successfully created
```

```

| reason: MachineCreationSucceeded
| status: "True"
| type: MachineCreation
| instanceId: i-0fdb85790d76d0c3f
| instanceState: stopped
| kind: AWSMachineProviderStatus

```

- ii. Remove the **providerID** field:

```

| providerID: aws:///us-east-1a/i-0fdb85790d76d0c3f

```

- iii. Change the **name** field to a new name.

It is recommended to keep the same base name as the old machine and change the ending number to the next available number. In this example, **clusternode-8qw5l-master-0** is changed to **clusternode-8qw5l-master-3**.

For example:

```

| apiVersion: machine.openshift.io/v1beta1
| kind: Machine
| metadata:
| ...
|   name: clusternode-8qw5l-master-3
| ...

```

- iv. Update the **selfLink** field to use the new machine name from the previous step.

```

| apiVersion: machine.openshift.io/v1beta1
| kind: Machine
| metadata:
| ...
|   selfLink: /apis/machine.openshift.io/v1beta1/namespaces/openshift-machine-
|             api/machines/clusternode-8qw5l-master-3
| ...

```

- d. Delete the machine of the unhealthy member:

```

| $ oc delete machine -n openshift-machine-api clusternode-8qw5l-master-0 ①

```

- ① Specify the name of the master machine for the unhealthy node.

- e. Verify that the machine was deleted:

```

| $ oc get machines -n openshift-machine-api -o wide

```

Example output

NAME NODE	PHASE PROVIDERID	TYPE STATE	REGION	ZONE	AGE
clusternode-8qw5l-master-1 3h37m		Running aws:///us-east-1b/i-096c349b700a19631	m4.xlarge	us-east-1	us-east-1b
clusternode-8qw5l-master-2		Running	m4.xlarge	us-east-1	us-east-1c

```
3h37m ip-10-0-164-97.ec2.internal aws:///us-east-1c/i-02626f1dba9ed5bba running
clustername-8qw5l-worker-us-east-1a-wbtgd Running m4.large us-east-1 us-east-1a
3h28m ip-10-0-129-226.ec2.internal aws:///us-east-1a/i-010ef6279b4662ced
running
clustername-8qw5l-worker-us-east-1b-lrdxb Running m4.large us-east-1 us-east-1b
3h28m ip-10-0-144-248.ec2.internal aws:///us-east-1b/i-0cb45ac45a166173b running
clustername-8qw5l-worker-us-east-1c-pkg26 Running m4.large us-east-1 us-east-1c
3h28m ip-10-0-170-181.ec2.internal aws:///us-east-1c/i-06861c00007751b0a
running
```

- f. Create the new machine using the **new-master-machine.yaml** file:

```
$ oc apply -f new-master-machine.yaml
```

- g. Verify that the new machine has been created:

```
$ oc get machines -n openshift-machine-api -o wide
```

Example output

NAME NODE	PHASE PROVIDERID	TYPE	REGION	ZONE	AGE STATE
clustername-8qw5l-master-1 3h37m ip-10-0-154-204.ec2.internal	Running aws:///us-east-1b/i-096c349b700a19631	m4.xlarge	us-east-1	us-east-1b	running
clustername-8qw5l-master-2 3h37m ip-10-0-164-97.ec2.internal	Running aws:///us-east-1c/i-02626f1dba9ed5bba	m4.xlarge	us-east-1	us-east-1c	running
clustername-8qw5l-master-3 85s ip-10-0-133-53.ec2.internal	Provisioning aws:///us-east-1a/i-015b0888fe17bc2c8	m4.xlarge	us-east-1	us-east-1a	running
① clustername-8qw5l-worker-us-east-1a-wbtgd 3h28m ip-10-0-129-226.ec2.internal running	Running aws:///us-east-1a/i-010ef6279b4662ced	m4.large	us-east-1	us-east-1a	
	Running aws:///us-east-1b/i-0cb45ac45a166173b	m4.large	us-east-1	us-east-1b	running
	Running aws:///us-east-1c/i-06861c00007751b0a	m4.large	us-east-1	us-east-1c	running

- ① The new machine, **clustername-8qw5l-master-3** is being created and is ready once the phase changes from **Provisioning** to **Running**.

It might take a few minutes for the new machine to be created. The etcd cluster Operator will automatically sync when the machine or node returns to a healthy state.

3. Verify that all etcd Pods are running properly:

In a terminal that has access to the cluster as a **cluster-admin** user, run the following command:

```
$ oc get pods -n openshift-etcd | grep etcd
```

Example output

etcd-ip-10-0-133-53.ec2.internal	3/3	Running	0	7m49s
etcd-ip-10-0-164-97.ec2.internal	3/3	Running	0	123m
etcd-ip-10-0-154-204.ec2.internal	3/3	Running	0	124m

2.4.2. Replacing an unhealthy etcd member whose etcd Pod is crashlooping

This procedure details the steps to replace an etcd member that is unhealthy because the etcd Pod is crashlooping.

Prerequisites

- You have identified the unhealthy etcd member.
- You have verified that the etcd Pod is crashlooping.
- You have access to the cluster as a user with the **cluster-admin** role.
- You have taken an etcd backup.



IMPORTANT

It is important to take an etcd backup before performing this procedure so that your cluster can be restored if you encounter any issues.

Procedure

1. Stop the crashlooping etcd Pod.
 - a. Debug the node that is crashlooping.
In a terminal that has access to the cluster as a **cluster-admin** user, run the following command:


```
$ oc debug node/ip-10-0-131-183.ec2.internal ①
```

① Replace this with the name of the unhealthy node.
 - b. Change your root directory to the host:


```
sh-4.2# chroot /host
```
 - c. Move the existing etcd Pod file out of the kubelet manifest directory:


```
sh-4.2# mv /etc/kubernetes/manifests/etcd-pod.yaml /var/lib/etcd-backup/
```
 - d. Move the etcd data directory to a different location:


```
sh-4.2# mv /var/lib/etcd/ /tmp
```

You can now exit the node shell.
2. Remove the unhealthy member.

- Choose a Pod that is *not* on the affected node.

In a terminal that has access to the cluster as a **cluster-admin** user, run the following command:

```
$ oc get pods -n openshift-etcd | grep etcd
```

Example output

etcd-ip-10-0-131-183.ec2.internal	2/3	Error	7	6h9m
etcd-ip-10-0-164-97.ec2.internal	3/3	Running	0	6h6m
etcd-ip-10-0-154-204.ec2.internal	3/3	Running	0	6h6m

- Connect to the running etcd container, passing in the name of a Pod that is not on the affected node.

In a terminal that has access to the cluster as a **cluster-admin** user, run the following command:

```
$ oc rsh -n openshift-etcd etcd-ip-10-0-154-204.ec2.internal
```

- View the member list:

```
sh-4.2# etcdctl member list -w table
```

Example output

ID ADDRS	STATUS	NAME	PEER ADDRS	CLIENT
62bcf33650a7170a	started	ip-10-0-131-183.ec2.internal	https://10.0.131.183:2380	
			https://10.0.131.183:2379	
b78e2856655bc2eb	started	ip-10-0-164-97.ec2.internal	https://10.0.164.97:2380	
			https://10.0.164.97:2379	
d022e10b498760d5	started	ip-10-0-154-204.ec2.internal	https://10.0.154.204:2380	
			https://10.0.154.204:2379	

- Remove the unhealthy etcd member by providing the ID to the **etcdctl member remove** command:

```
sh-4.2# etcdctl member remove 62bcf33650a7170a
```

Example output

```
Member 62bcf33650a7170a removed from cluster ead669ce1fbfb346
```

- View the member list again and verify that the member was removed:

```
sh-4.2# etcdctl member list -w table
```

Example output

ID ADDRS	STATUS	NAME	PEER ADDRS	CLIENT
b78e2856655bc2eb started ip-10-0-164-97.ec2.internal https://10.0.164.97:2380 https://10.0.164.97:2379				
d022e10b498760d5 started ip-10-0-154-204.ec2.internal https://10.0.154.204:2380 https://10.0.154.204:2379				

You can now exit the node shell.

3. Force etcd redeployment.

In a terminal that has access to the cluster as a **cluster-admin** user, run the following command:

```
$ oc patch etcd cluster -p='{"spec": {"forceRedeploymentReason": "single-master-recovery-  
"$( date --rfc-3339=ns )""}}' --type=merge ①
```

- ① The **forceRedeploymentReason** value must be unique, which is why a timestamp is appended.

When the etcd cluster Operator performs a redeployment, it ensures that all master nodes have a functioning etcd Pod.

4. Verify that the new member is available and healthy.

- Connect to the running etcd container again.

In a terminal that has access to the cluster as a cluster-admin user, run the following command:

```
$ oc rsh -n openshift-etcd etcd-ip-10-0-154-204.ec2.internal
```

- Verify that all members are healthy:

```
sh-4.2# etcdctl endpoint health --cluster
```

Example output

```
https://10.0.131.183:2379 is healthy: successfully committed proposal: took =  
16.671434ms  
https://10.0.154.204:2379 is healthy: successfully committed proposal: took =  
16.698331ms  
https://10.0.164.97:2379 is healthy: successfully committed proposal: took =  
16.621645ms
```

CHAPTER 3. SHUTTING DOWN THE CLUSTER GRACEFULLY

This document describes the process to gracefully shut down your cluster. You might need to temporarily shut down your cluster for maintenance reasons, or to save on resource costs.

3.1. PREREQUISITES

- Take an [etcd backup](#) prior to shutting down the cluster.

3.2. SHUTTING DOWN THE CLUSTER

You can shut down your cluster in a graceful manner so that it can be restarted at a later date.

Prerequisites

- You have access to the cluster as a user with the **cluster-admin** role.
- You have taken an etcd backup.



IMPORTANT

It is important to take an etcd backup before performing this procedure so that your cluster can be restored if you encounter any issues when restarting the cluster.

Procedure

1. Shut down all of the nodes in the cluster. You can do this from your cloud provider's web console, or you can use the below commands:

- a. Obtain the list of nodes:

```
$ nodes=$(oc get nodes -o name)
```

- b. Shut down all of the nodes:

```
$ for node in ${nodes[@]}\ndo\n    echo "==== Shut down $node ====\n    ssh core@$node sudo shutdown -h 1\ndone
```

Shutting down the nodes using one of these methods allows pods to terminate gracefully, which reduces the chance for data corruption.



NOTE

It is not necessary to drain master nodes of the standard Pods that ship with OpenShift Container Platform prior to shutdown.

Cluster administrators are responsible for ensuring a clean restart of their own workloads after the cluster is restarted. If you drained master nodes prior to shutdown because of custom workloads, you must mark the master nodes as schedulable before the cluster will be functional again after restart.

2. Shut off any cluster dependencies that are no longer needed, such as external storage or an LDAP server. Be sure to consult your vendor's documentation before doing so.

Additional resources

- [Restarting the cluster gracefully](#)

CHAPTER 4. RESTARTING THE CLUSTER GRACEFULLY

This document describes the process to restart your cluster after a graceful shutdown.

Even though the cluster is expected to be functional after the restart, the cluster might not recover due to unexpected conditions, for example:

- etcd data corruption during shutdown
- Node failure due to hardware
- Network connectivity issues

If your cluster fails to recover, follow the steps to [restore to a previous cluster state](#).

4.1. PREREQUISITES

- You have [gracefully shut down your cluster](#).

4.2. RESTARTING THE CLUSTER

You can restart your cluster after it has been shut down gracefully.

Prerequisites

- You have access to the cluster as a user with the **cluster-admin** role.
- This procedure assumes that you gracefully shut down the cluster.

Procedure

1. Power on any cluster dependencies, such as external storage or an LDAP server.
2. Start all cluster machines.
Use the appropriate method for your cloud environment to start the machines, for example, from your cloud provider's web console.
Wait approximately 10 minutes before continuing to check the status of master nodes.
3. Verify that all master nodes are ready.

```
$ oc get nodes -l node-role.kubernetes.io/master
```

The master nodes are ready if the status is **Ready**, as shown in the following output:

NAME	STATUS	ROLES	AGE	VERSION
ip-10-0-168-251.ec2.internal	Ready	master	75m	v1.18.3
ip-10-0-170-223.ec2.internal	Ready	master	75m	v1.18.3
ip-10-0-211-16.ec2.internal	Ready	master	75m	v1.18.3

4. If the master nodes are *not* ready, then check whether there are any pending certificate signing requests that must be approved.
 - a. Get the list of current CSRs:

```
$ oc get csr
```

- b. Review the details of a CSR to verify that it is valid:

```
$ oc describe csr <csr_name> 1
```

1 <csr_name> is the name of a CSR from the list of current CSRs.

- c. Approve each valid CSR:

```
$ oc adm certificate approve <csr_name>
```

5. After the master nodes are ready, verify that all worker nodes are ready.

```
$ oc get nodes -l node-role.kubernetes.io/worker
```

The worker nodes are ready if the status is **Ready**, as shown in the following output:

NAME	STATUS	ROLES	AGE	VERSION
ip-10-0-179-95.ec2.internal	Ready	worker	64m	v1.18.3
ip-10-0-182-134.ec2.internal	Ready	worker	64m	v1.18.3
ip-10-0-250-100.ec2.internal	Ready	worker	64m	v1.18.3

6. If the worker nodes are *not* ready, then check whether there are any pending certificate signing requests that must be approved.

- a. Get the list of current CSRs:

```
$ oc get csr
```

- b. Review the details of a CSR to verify that it is valid:

```
$ oc describe csr <csr_name> 1
```

1 <csr_name> is the name of a CSR from the list of current CSRs.

- c. Approve each valid CSR:

```
$ oc adm certificate approve <csr_name>
```

7. Verify that the cluster started properly.

- a. Check that there are no degraded cluster Operators.

```
$ oc get clusteroperators
```

Check that there are no cluster Operators with the **DEGRADED** condition set to **True**.

NAME	VERSION	AVAILABLE	PROGRESSING	DEGRADED
SINCE				
authentication	4.5.0	True	False	False
				59m

cloud-credential	4.5.0	True	False	False	85m
cluster-autoscaler	4.5.0	True	False	False	73m
config-operator	4.5.0	True	False	False	73m
console	4.5.0	True	False	False	62m
csi-snapshot-controller	4.5.0	True	False	False	66m
dns	4.5.0	True	False	False	76m
etcd	4.5.0	True	False	False	76m
...					

- b. Check that all nodes are in the ready state:

```
$ oc get nodes
```

Check that the status for all nodes is **Ready**.

NAME	STATUS	ROLES	AGE	VERSION
ip-10-0-168-251.ec2.internal	Ready	master	82m	v1.18.3
ip-10-0-170-223.ec2.internal	Ready	master	82m	v1.18.3
ip-10-0-179-95.ec2.internal	Ready	worker	70m	v1.18.3
ip-10-0-182-134.ec2.internal	Ready	worker	70m	v1.18.3
ip-10-0-211-16.ec2.internal	Ready	master	82m	v1.18.3
ip-10-0-250-100.ec2.internal	Ready	worker	69m	v1.18.3

If the cluster did not start properly, you might need to restore your cluster using an etcd backup.

Additional resources

- See [Restoring to a previous cluster state](#) for how to use an etcd backup to restore if your cluster failed to recover after restarting.

CHAPTER 5. DISASTER RECOVERY

5.1. ABOUT DISASTER RECOVERY

The disaster recovery documentation provides information for administrators on how to recover from several disaster situations that might occur with their OpenShift Container Platform cluster. As an administrator, you might need to follow one or more of the following procedures in order to return your cluster to a working state.

Restoring to a previous cluster state

This solution handles situations where you want to restore your cluster to a previous state, for example, if an administrator deletes something critical. This also includes situations where you have lost the majority of your master hosts, leading to etcd quorum loss and the cluster going offline. As long as you have taken an etcd backup, you can follow this procedure to restore your cluster to a previous state.

If applicable, you might also need to [recover from expired control plane certificates](#).



NOTE

If you have a majority of your masters still available and have an etcd quorum, then follow the procedure to [replace a single unhealthy etcd member](#).

Recovering from expired control plane certificates

This solution handles situations where your control plane certificates have expired. For example, if you shut down your cluster before the first certificate rotation, which occurs 24 hours after installation, your certificates will not be rotated and will expire. You can follow this procedure to recover from expired control plane certificates.

5.2. RECOVERING FROM LOST MASTER HOSTS

As of OpenShift Container Platform 4.4, follow the procedure to [restore to a previous cluster state](#) in order to recover from lost master hosts.



NOTE

If you have a majority of your masters still available and have an etcd quorum, then follow the procedure to [replace a single unhealthy etcd member](#).

5.3. RESTORING TO A PREVIOUS CLUSTER STATE

To restore the cluster to a previous state, you must have previously [backed up etcd data](#) by creating a snapshot. You will use this snapshot to restore the cluster state.

5.3.1. Restoring to a previous cluster state

You can use a saved etcd backup to restore back to a previous cluster state. You use the etcd backup to restore a single control plane host. Then the etcd cluster Operator handles scaling to the remaining master hosts.

Prerequisites

- Access to the cluster as a user with the **cluster-admin** role.
- SSH access to master hosts.
- A backup directory containing both the etcd snapshot and the resources for the static Pods, which were from the same backup. The file names in the directory must be in the following formats: **snapshot_<datetimestamp>.db** and **static_kuberresources_<datetimestamp>.tar.gz**.

Procedure

1. Select a control plane host to use as the recovery host. This is the host that you will run the restore operation on.
2. Establish SSH connectivity to each of the control plane nodes, including the recovery host. The Kubernetes API server becomes inaccessible after the restore process starts, so you cannot access the control plane nodes. For this reason, it is recommended to establish SSH connectivity to each control plane host in a separate terminal.



IMPORTANT

If you do not complete this step, you will not be able to access the master hosts to complete the restore procedure, and you will be unable to recover your cluster from this state.

3. Copy the etcd backup directory to the recovery control plane host.
This procedure assumes that you copied the **backup** directory containing the etcd snapshot and the resources for the static Pods to the **/home/core/** directory of your recovery control plane host.
4. Stop the static Pods on all other control plane nodes.



NOTE

It is not required to manually stop the Pods on the recovery host. The recovery script will stop the Pods on the recovery host.

- a. Access a control plane host that is not the recovery host.

- b. Move the existing etcd Pod file out of the kubelet manifest directory:

```
[core@ip-10-0-154-194 ~]$ sudo mv /etc/kubernetes/manifests/etcd-pod.yaml /tmp
```

- c. Verify that the etcd Pods are stopped.

```
[core@ip-10-0-154-194 ~]$ sudo crictl ps | grep etcd
```

The output of this command should be empty. If it is not empty, wait a few minutes and check again.

- d. Move the existing Kubernetes API server Pod file out of the kubelet manifest directory:

```
[core@ip-10-0-154-194 ~]$ sudo mv /etc/kubernetes/manifests/kube-apiserver-pod.yaml /tmp
```

- - e. Verify that the Kubernetes API server Pods are stopped.

```
[core@ip-10-0-154-194 ~]$ sudo crictl ps | grep kube-apiserver
```

The output of this command should be empty. If it is not empty, wait a few minutes and check again.

- f. Move the etcd data directory to a different location:

```
[core@ip-10-0-154-194 ~]$ sudo mv /var/lib/etcd/ /tmp
```

- g. Repeat this step on each of the other master hosts that is not the recovery host.

5. Access the recovery control plane host.
6. If the cluster-wide proxy is enabled, be sure that you have exported the **NO_PROXY**, **HTTP_PROXY**, and **HTTPS_PROXY** environment variables.

TIP

You can check whether the proxy is enabled by reviewing the output of **oc get proxy cluster -o yaml**. The proxy is enabled if the **httpProxy**, **httpsProxy**, and **noProxy** fields have values set.

7. Run the restore script on the recovery control plane host and pass in the path to the etcd backup directory:

```
[core@ip-10-0-143-125 ~]$ sudo -E /usr/local/bin/cluster-restore.sh /home/core/backup
```

Example script output

```
...stopping kube-scheduler-pod.yaml  
...stopping kube-controller-manager-pod.yaml  
...stopping etcd-pod.yaml  
...stopping kube-apiserver-pod.yaml  
Waiting for container etcd to stop  
.complete  
Waiting for container etcdctl to stop  
.....complete  
Waiting for container etcd-metrics to stop  
complete  
Waiting for container kube-controller-manager to stop  
complete  
Waiting for container kube-apiserver to stop  
.....complete  
Waiting for container kube-scheduler to stop  
complete  
Moving etcd data-dir /var/lib/etcd/member to /var/lib/etcd-backup  
starting restore-etcd static pod  
starting kube-apiserver-pod.yaml  
static-pod-resources/kube-apiserver-pod-7/kube-apiserver-pod.yaml  
starting kube-controller-manager-pod.yaml
```

```
static-pod-resources/kube-controller-manager-pod-7/kube-controller-manager-pod.yaml
starting kube-scheduler-pod.yaml
static-pod-resources/kube-scheduler-pod-8/kube-scheduler-pod.yaml
```

8. Restart the kubelet service on all master hosts.

- a. From the recovery host, run the following command:

```
[core@ip-10-0-143-125 ~]$ sudo systemctl restart kubelet.service
```

- b. Repeat this step on all other master hosts.

9. Verify that the single member control plane has started successfully.

- a. From the recovery host, verify that the etcd container is running.

```
[core@ip-10-0-143-125 ~]$ sudo crictl ps | grep etcd
```

Example output

```
3ad41b7908e32
36f86e2eeaaffe662df0d21041eb22b8198e0e58abeeae8c743c3e6e977e8009
About a minute ago  Running      etcd          0
7c05f8af362f0
```

- b. From the recovery host, verify that the etcd Pod is running.

```
[core@ip-10-0-143-125 ~]$ oc get pods -n openshift-etcd | grep etcd
```

Example output

NAME	READY	STATUS	RESTARTS	AGE
etcd-ip-10-0-143-125.ec2.internal	1/1	Running	1	2m47s

If the status is **Pending**, or the output lists more than one running etcd Pod, wait a few minutes and check again.

10. Force etcd redeployment.

In a terminal that has access to the cluster as a **cluster-admin** user, run the following command:

```
$ oc patch etcd cluster -p='{"spec": {"forceRedeploymentReason": "recovery-"'$( date --rfc-3339=ns )'"}}' --type=merge ①
```

- 1 The **forceRedeploymentReason** value must be unique, which is why a timestamp is appended.

When the etcd cluster Operator performs a redeployment, the existing nodes are started with new Pods similar to the initial bootstrap scale up.

11. Verify all nodes are updated to the latest revision.

In a terminal that has access to the cluster as a **cluster-admin** user, run the following command:

```
$ oc get etcd -o=jsonpath='{range .items[0].status.conditions[?(@.type=="NodeInstallerProgressing")]}{.reason}{"\n"}{.message}{"\n"}'
```

Review the **NodeInstallerProgressing** status condition for etcd to verify that all nodes are at the latest revision. The output shows **AllNodesAtLatestRevision** upon successful update:

```
AllNodesAtLatestRevision
3 nodes are at revision 3
```

If the output shows a message such as **2 nodes are at revision 3; 1 nodes are at revision 4**, this means that the update is still in progress. Wait a few minutes and try again.

12. After etcd is redeployed, force new rollouts for the control plane. The Kubernetes API server will reinstall itself on the other nodes because the kubelet is connected to API servers using an internal load balancer.

In a terminal that has access to the cluster as a **cluster-admin** user, run the following commands.

- a. Update the **kubeapiserver**:

```
$ oc patch kubeapiserver cluster -p='{"spec": {"forceRedeploymentReason": "recovery-""$( date --rfc-3339=ns )"""}' --type=merge
```

Verify all nodes are updated to the latest revision.

```
$ oc get kubeapiserver -o=jsonpath='{range .items[0].status.conditions[?(@.type=="NodeInstallerProgressing")]}{.reason}{"\n"}{.message}{"\n"}'
```

Review the **NodeInstallerProgressing** status condition to verify that all nodes are at the latest revision. The output shows **AllNodesAtLatestRevision** upon successful update:

```
AllNodesAtLatestRevision
3 nodes are at revision 3
```

- b. Update the **kubecontrollermanager**:

```
$ oc patch kubecontrollermanager cluster -p='{"spec": {"forceRedeploymentReason": "recovery-""$( date --rfc-3339=ns )"""}' --type=merge
```

Verify all nodes are updated to the latest revision.

```
$ oc get kubecontrollermanager -o=jsonpath='{range .items[0].status.conditions[?(@.type=="NodeInstallerProgressing")]}{.reason}{"\n"}{.message}{"\n"}'
```

Review the **NodeInstallerProgressing** status condition to verify that all nodes are at the latest revision. The output shows **AllNodesAtLatestRevision** upon successful update:

```
AllNodesAtLatestRevision
3 nodes are at revision 3
```

- c. Update the **kubescheduler**:

```
$ oc patch kubescheduler cluster -p='{"spec": {"forceRedeploymentReason": "recovery-"
"$( date --rfc-3339=ns )"""}' --type=merge
```

Verify all nodes are updated to the latest revision.

```
$ oc get kubescheduler -o=jsonpath='{range .items[0].status.conditions[?
(@.type=="NodeInstallerProgressing")]}{.reason}{"\n"}{.message}{"\n"}'
```

Review the **NodeInstallerProgressing** status condition to verify that all nodes are at the latest revision. The output shows **AllNodesAtLatestRevision** upon successful update:

```
AllNodesAtLatestRevision
3 nodes are at revision 3
```

13. Verify that all master hosts have started and joined the cluster.

In a terminal that has access to the cluster as a **cluster-admin** user, run the following command:

```
$ oc get pods -n openshift-etcd | grep etcd
```

Example output

etcd-ip-10-0-143-125.ec2.internal	2/2	Running	0	9h
etcd-ip-10-0-154-194.ec2.internal	2/2	Running	0	9h
etcd-ip-10-0-173-171.ec2.internal	2/2	Running	0	9h

Note that it might take several minutes after completing this procedure for all services to be restored. For example, authentication by using **oc login** might not immediately work until the OAuth server Pods are restarted.

5.4. RECOVERING FROM EXPIRED CONTROL PLANE CERTIFICATES

5.4.1. Recovering from expired control plane certificates

As of OpenShift Container Platform 4.4.8, the cluster can automatically recover from expired control plane certificates. You no longer need to perform the manual steps that were required in previous versions.

The exception is that you must manually approve the pending **node-bootstrapper** certificate signing requests (CSRs) to recover kubelet certificates.

Use the following steps to approve the pending **node-bootstrapper** CSRs.

Procedure

1. Get the list of current CSRs:

```
$ oc get csr
```

2. Review the details of a CSR to verify that it is valid:

```
$ oc describe csr <csr_name> ①
```

- 1 <csr_name> is the name of a CSR from the list of current CSRs.

3. Approve each valid **node-bootstrapper** CSR:

```
$ oc adm certificate approve <csr_name>
```